

week6 小红书营销渠道效果预测分析

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
```

1.数据查看

```
In [2]: data=pd.read_csv('C:\\Users\\mac\\Desktop\\数据分析班\\week6\\小红书数据.csv')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29452 entries, 0 to 29451
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   revenue                29452 non-null  float64
1   gender                 17723 non-null  float64
2   age                    17723 non-null  float64
3   engaged_last_30        17723 non-null  float64
4   lifecycle              29452 non-null  object
5   days_since_last_order  29452 non-null  float64
6   previous_order_amount  29452 non-null  float64
7   3rd_party_stores       29452 non-null  int64
dtypes: float64(6), int64(1), object(1)
memory usage: 1.8+ MB
```

gender, age和engaged_last_30有缺失值
gender和engaged_last_30的数据类型错误

```
In [3]: data.isnull().sum()
```

```
Out[3]: revenue                0
gender                 11729
age                    11729
engaged_last_30        11729
lifecycle              0
days_since_last_order  0
previous_order_amount  0
3rd_party_stores       0
dtype: int64
```

缺失值比例

```
In [4]: data.isnull().sum()/data.shape[0]
```

```
Out[4]: revenue                0.000000
gender                 0.398241
age                    0.398241
engaged_last_30        0.398241
lifecycle              0.000000
days_since_last_order  0.000000
previous_order_amount  0.000000
```

3rd_party_stores

0.000000

dtype: float64

比例接近40%，不可直接删除，必须只能填充

In [5]:

```
data.head()
```

Out[5]:

	revenue	gender	age	engaged_last_30	lifecycle	days_since_last_order	previous_order_amount
0	72.98	0.0	43.0	0.0	B	4.26	2343.870
1	200.99	0.0	34.0	0.0	A	0.94	8539.872
2	69.98	0.0	16.0	0.0	C	4.29	1687.646
3	649.99	NaN	NaN	NaN	C	14.90	3498.846
4	83.59	NaN	NaN	NaN	C	21.13	3968.490

In [6]:

```
data.describe()
```

Out[6]:

	revenue	gender	age	engaged_last_30	days_since_last_order	previous_c
count	29452.000000	17723.000000	17723.000000	17723.000000	29452.000000	
mean	397.071515	0.298200	29.419286	0.073069	7.711348	
std	959.755615	0.457481	9.213604	0.260257	6.489289	
min	0.020000	0.000000	14.000000	0.000000	0.130000	
25%	74.970000	0.000000	21.000000	0.000000	2.190000	
50%	175.980000	0.000000	29.000000	0.000000	5.970000	
75%	498.772500	1.000000	37.000000	0.000000	11.740000	
max	103466.100000	1.000000	45.000000	1.000000	23.710000	

发现：revenue和previous_order_amount标准差太大，回归分析时需要偏离较大的值


2收入EDA

In [7]:

```
data.revenue.plot(kind="hist")
```


Out[7]:

```
<AxesSubplot:ylabel='Frequency'>
```




```
In [8]: data[data.revenue<10000]['revenue'].plot(kind="hist")
```

```
Out[8]: <AxesSubplot:ylabel='Frequency'>
```




```
In [9]: data[data.revenue<4000]['revenue'].plot(kind="hist")
```

```
Out[9]: <AxesSubplot:ylabel='Frequency'>
```



```
In [10]: data[data.revenue<2000]['revenue'].plot(kind="hist")
```

```
Out[10]: <AxesSubplot:ylabel='Frequency'>
```



分布在0~1000的人数占绝大多数

```
In [11]: data.revenue.describe()
```

```
Out[11]: count      29452.000000
         mean        397.071515
         std         959.755615
         min          0.020000
         25%         74.970000
         50%        175.980000
         75%        498.772500
         max       103466.100000
         Name: revenue, dtype: float64
```

```
In [12]: diff=data.revenue.describe()["75%"]-data.revenue.describe()["25%"]
         net_max=data.revenue.describe()["75%"]+1.5*diff
         net_max
```

```
Out[12]: 1134.47625
```

故可以认为大于1134.47625的均为revenue的离群值

```
In [13]: 1-data[data.revenue<=1134.47625].shape[0]/data.shape[0]
```

```
Out[13]: 0.07357734619041156
```

离群值占比7%左右，删去是可以接受的

3.相关可视化分析

```
In [14]: data.corr()["revenue"]
```


```
Out[14]: revenue      1.000000
         gender      0.014944
         age         0.006263
         engaged_last_30 0.080031
         days_since_last_order 0.036754
         previous_order_amount 0.168186
         3rd_party_stores -0.026102
         Name: revenue, dtype: float64
```

由此可得, `previous_order_amount` 与 `revenue` 的相关性最强, 其次是 `engaged_last_30` 和 `days_since_last_order`, 但都很低

```
In [15]: data=data[data.revenue<=1134.47625]
```


```
In [16]: sns.regplot(x='previous_order_amount',y='revenue',data=data)
```

```
Out[16]: <AxesSubplot:xlabel='previous_order_amount', ylabel='revenue'>
```




```
In [17]: sns.regplot(x='engaged_last_30',y='revenue',data=data)
```

```
Out[17]: <AxesSubplot:xlabel='engaged_last_30', ylabel='revenue'>
```




```
In [18]: sns.regplot(x=" days_since_last_order ",y='revenue',data=data)
```

```
Out[18]: <AxesSubplot:xlabel=' days_since_last_order ', ylabel='revenue'>
```




```
In [19]: sns.barplot(x="gender", y="revenue", data=data)
```

```
Out[19]: <AxesSubplot:xlabel='gender', ylabel='revenue'>
```



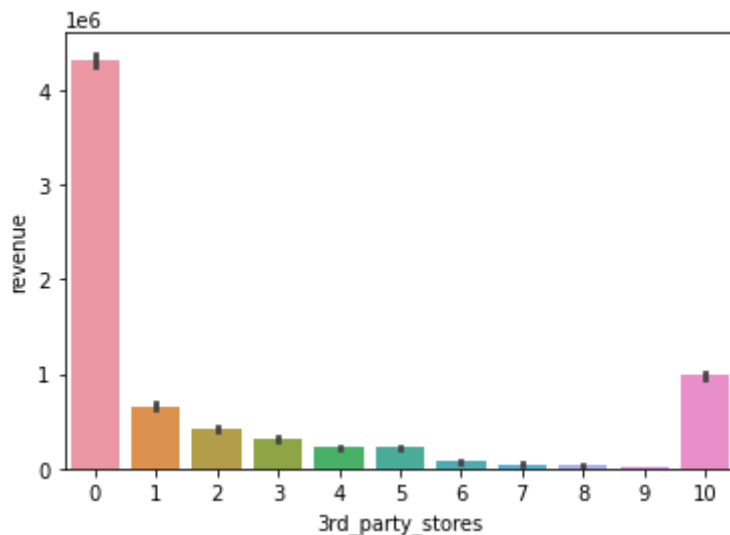
```
In [20]: sns.barplot(x="3rd_party_stores", y="revenue", data=data)
```

```
Out[20]: <AxesSubplot:xlabel='3rd_party_stores', ylabel='revenue'>
```



```
In [21]: sns.barplot(x="3rd_party_stores", y="revenue", estimator=sum, data=data)
```

```
Out[21]: <AxesSubplot:xlabel='3rd_party_stores', ylabel='revenue'>
```



4.数据处理

4.1 revenue数据清洗

```
In [22]: df=data.copy()
```

```
In [23]: df=df[df.revenue<=1134.47625]
```

4.2缺失值填充


4.2.1性别

```
In [24]: df.gender.unique()
```

```
Out[24]: array([ 0., nan,  1.])
```


```
In [25]: df['gender'].value_counts().plot(kind="bar")
```

```
Out[25]: <AxesSubplot:>
```



```
In [26]: df['gender'].value_counts(dropna=False).plot(kind="bar")
```

```
Out[26]: <AxesSubplot:>
```




```
In [27]: df['gender']=df['gender'].fillna('unknown')
```

4.2.2年龄

```
In [28]: df['age'].plot(kind="hist",bins=100)
```

```
Out[28]: <AxesSubplot:ylabel='Frequency'>
```



```
In [29]: df['age']=df['age'].fillna(df['age'].mean())
```

4.2.3过去三十天参与的活动

```
In [30]: df['engaged_last_30']=df['engaged_last_30'].fillna(0.0)
```

4.2.4更改数据类型

```
In [31]: df['gender']=df['gender'].astype(str)
```



```
In [32]: df['engaged_last_30']=df['engaged_last_30'].astype(str)
```

4.2.5添加哑变量

```
In [33]: df=pd.get_dummies(df,drop_first=True)
```

```
In [34]: df.head()
```

```
Out[34]:
```

	revenue	age	days_since_last_order	previous_order_amount	3rd_party_stores	gender_1.0	gender_2.0
0	72.98	43.000000	4.26	2343.870	0	0	0
1	200.99	34.000000	0.94	8539.872	0	0	0
2	69.98	16.000000	4.29	1687.646	1	0	0
3	649.99	29.402114	14.90	3498.846	0	0	0
4	83.59	29.402114	21.13	3968.490	4	0	0

5.使用Python建立线性回归模型

```
In [35]: from sklearn.linear_model import LinearRegression as lrg
```

```
In [60]: x=df[['previous_order_amount'," days_since_last_order "]]
y=df['revenue']
```

```
In [61]: model=lrg()
model.fit(x,y)
```

```
Out[61]: LinearRegression()
```

```
In [62]: model.coef_
```

```
Out[62]: array([0.01842654, 5.00525826])
```

```
In [63]: model.intercept_
```

```
Out[63]: 189.21781583073073
```

5.1 模型评估

```
In [65]: socre=model.score(x,y)
socre
```

```
Out[65]: 0.03505897168369887
```

5.1预测用户的消费金额变化

```
In [66]: y_pred=model.predict(x)
```

```
In [68]: error=y_pred-y
         error
```

```
Out[68]: 0      180.749624
         1      150.293029
         2      171.807846
         3     -321.722220
         4      284.514452
         ...
        29447    160.778538
        29448    212.411158
        29449    177.689914
        29450    266.004564
        29451   -478.274606
        Name: revenue, Length: 27285, dtype: float64
```

```
In [69]: rmse=(error**2).mean()**.5
         rmse
```

```
Out[69]: 263.2192954187874
```

```
In [70]: mae=abs(error).mean()
         mae
```

```
Out[70]: 208.21575378676076
```

```
In [71]: from statsmodels.formula.api import ols
         model_ols=ols('y~x', df).fit()
         print(model_ols.summary())
```

OLS Regression Results

Dep. Variable:	y	R-squared:	0.035			
Model:	OLS	Adj. R-squared:	0.035			
Method:	Least Squares	F-statistic:	495.6			
Date:	Tue, 15 Dec 2020	Prob (F-statistic):	3.76e-212			
Time:	20:53:22	Log-Likelihood:	-1.9077e+05			
No. Observations:	27285	AIC:	3.816e+05			
Df Residuals:	27282	BIC:	3.816e+05			
Df Model:	2					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
Intercept	189.2178	3.030	62.454	0.000	183.279	195.156
x[0]	0.0184	0.001	25.806	0.000	0.017	0.020
x[1]	5.0053	0.249	20.132	0.000	4.518	5.493
=====						
Omnibus:	5183.316	Durbin-Watson:	1.992			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	8675.546			
Skew:	1.302	Prob(JB):	0.00			
Kurtosis:	3.920	Cond. No.	6.02e+03			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 6.02e+03. This might indicate that there are strong multicollinearity or other numerical problems.